

HPC Cloud @ SURFsara

Design a Parallel Application



Natalie Danezi <anatoli.danezi@surfsara.nl>

Outline

- Parallel application:
 - What is the need? ... Laptop is not enough
 - How to make it parallel? ... Scalability
- Running a parallel application on the HPC cloud
- Hands-on Part B: demo



SURFsara @SURFsara_NL · Jun 22

Why do I Need SURFsara Cloud Facility in My Research? #abl #cfd slideshare.net /aliabbasicivil ... via @SlideShare

Highly demanding applications for **compute** and **storage** resources to:


- Analyse larger computational domains
- Analyse larger volume of data
- Achieve higher accuracy
- Retrieve the results faster

Scalability: system & software


- System architecture:
 - Distributed memory systems
 - Shared memory systems
- Software optimisation:
 - Dig in your code, e.g. MPI, OpenMP
 - Out of the box tools, e.g. XBeach, Delft3D, D-Flow FM

System architecture

- Distributed memory systems



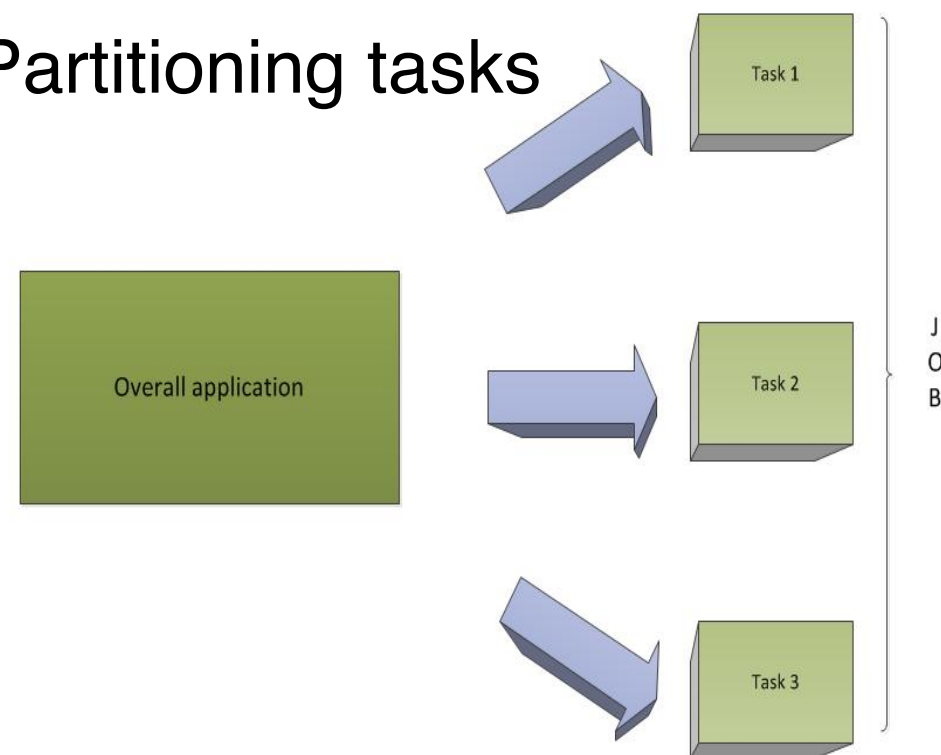
- Shared memory systems



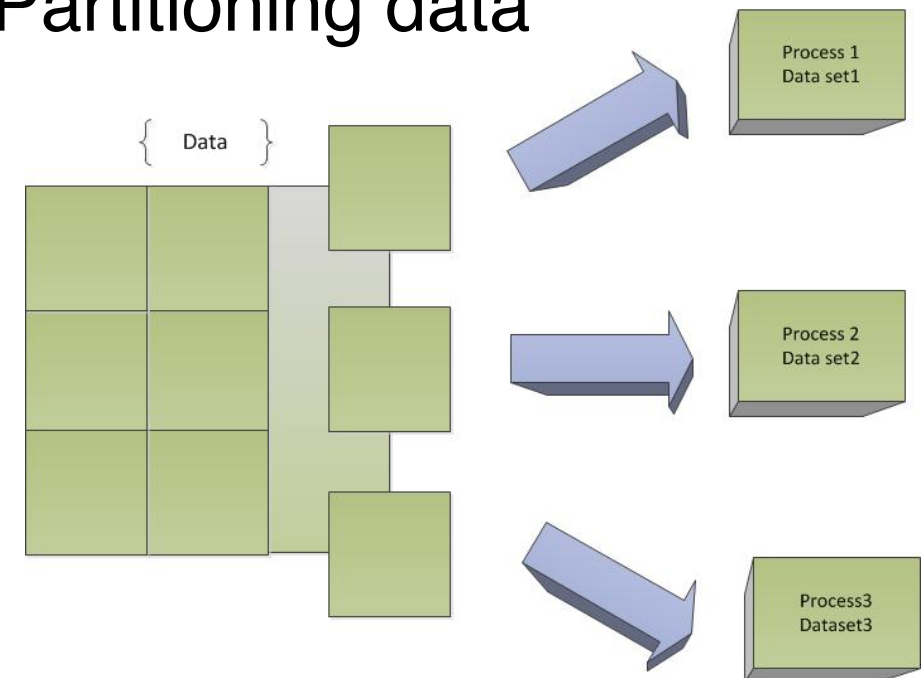
Software optimisation 1/2

- Dig in your code
 - MPI, OpenMP - advanced programming skills
 - Optimisation techniques - partitioning tasks / data

Partitioning tasks




Partitioning data



Software optimisation 2/2

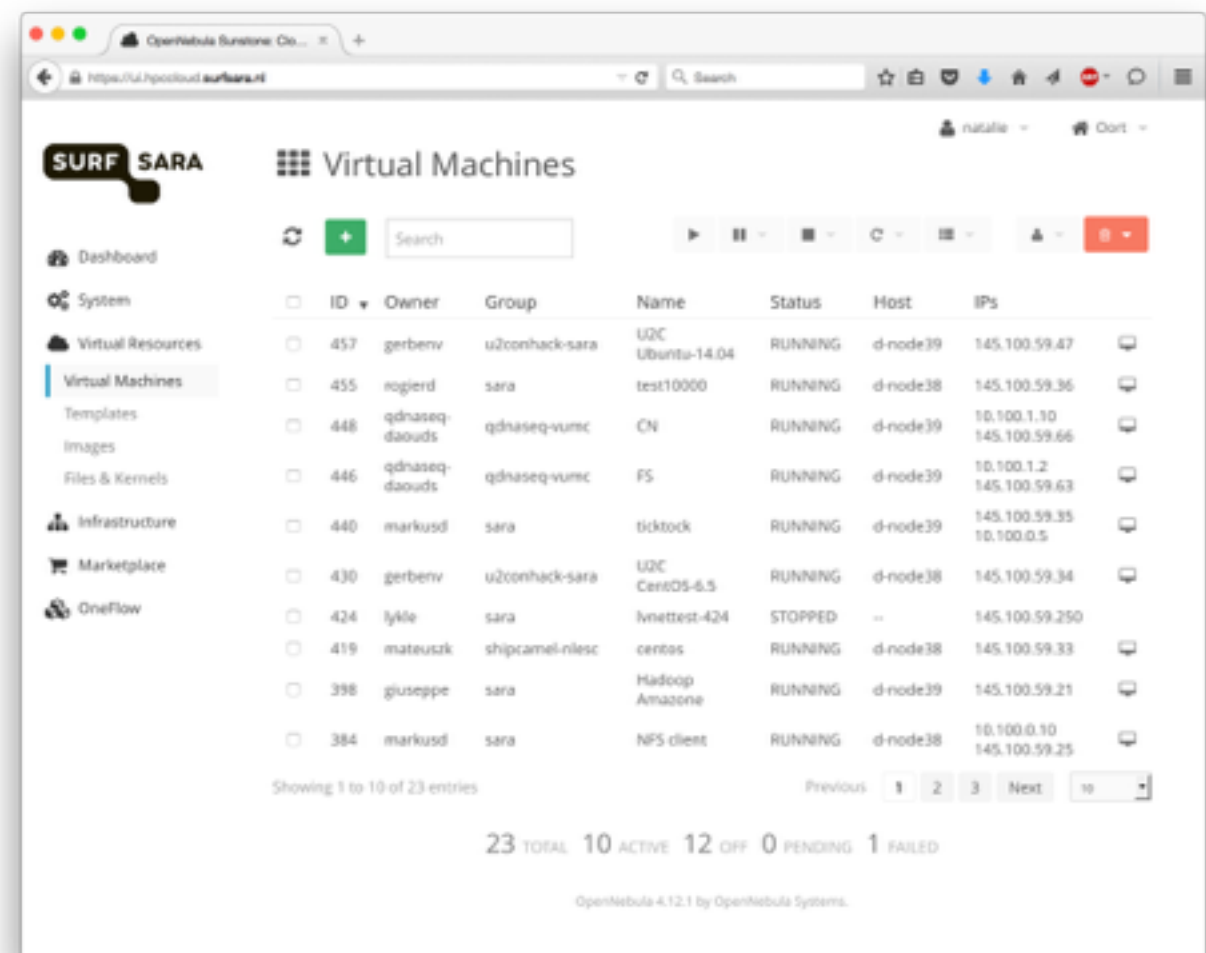
- Out of the box tools




Open  FOAM

Running a parallel application on the HPC cloud


- Access the UI
- Prepare OS Image, e.g. XBeach
- Create the Template
- Instantiate VM(s)
 - single-VM, single-core
 - single-VM, multi-core
 - multi-VM, multi-core
(master/workers cluster)



Scaling up vs. Scaling out



Scale up




Scale out



What makes HPC difficult?

- Scaling up
 - Amdahl's law
- Scaling out
 - More machines - more problems
 - Machines fail
 - Networks fail
 - Heterogeneous hardware
 - Latencies
 - Data locality
- Key questions...



Key questions

Define your specific requirements:

- Number of cores, Memory, Wall-clock time, data volumes
- Pipelines: what is shared in tasks, what needs to be communicated
- Data locality (private or shared)
- Wall-clock vs. CPU time
- Operating system, software (licensing programs) and databases
- Network interfaces (private or public) - virtual cluster
- Balance effort: system / software scalability



Hands-on Part B

UI: <https://ui.hpccloud.surfsara.nl>
Project: wshop-uihe
Username: wshop-sxz
Password: hydro@logyxz

- Continue from:
<http://doc.hpccloud.surfsara.nl/wshop-uihe-part-B-2015-Dec-11>
- Work in *pairs*
 - each with your *own credentials* on your own laptop
- Follow the instructions at your *own pace*
- Call *Natalie and Ander*
 - for a) doubts; b) when feeling stuck; c) “food for brain” hints